

# A Practical Guide to Running Experiments in RL

Subho, April 2023

# Overview

- A Typical Workflow for RL Experiments.
- Automating various aspects of the **experiment pipeline**.
  - Experiment Tracking (**Weights and Biases**)
  - Logging and Configs (**Hydra**)
  - Running hyper parameter sweeps (**Submitit plugin**)
- Reproducible experiments using Singularity containers.

# Managing Experiments can be Cumbersome!

- Typical Empirical RL workflow:
  - Choose a simple problem setting that tests a specific property/answers a specific research question.
  - Identify set of baselines/approaches (**experiment**)
  - Repeat each experiment with multiple seeds, parameters (**run**)
  - Aggregate runs and compare results from experiments (**result**)

# Managing Experiments can be **Cumbersome!**

**This can be tricky to do in practice:**

As the number of experiments grows:

- Organizing data becomes cumbersome.
- Managing multiple experiment configurations can become messy.
- Scheduling sweep jobs in CC is not straightforward.
  - Improper scheduling strategies affect priority of jobs!
  - Training environments are difficult to replicate across multiple CC clusters.



A good experiment workflow **must** make it  
**easy** to iterate over multiple ideas

# My Current Workflow

1. Write a **Hydra** config file that specifies my experiment (combination of runs)
2. Schedule runs in Compute Canada using **submitit** (one run = one job).
3. Retrieve run data from **Wandb** Database using **Jupyter** Notebook.
4. Generate Results using **Matplotlib** + **Pandas**.




# Experiment Tracking

**Available tools: Weights and Biases, MLFlow, Comet.ml**

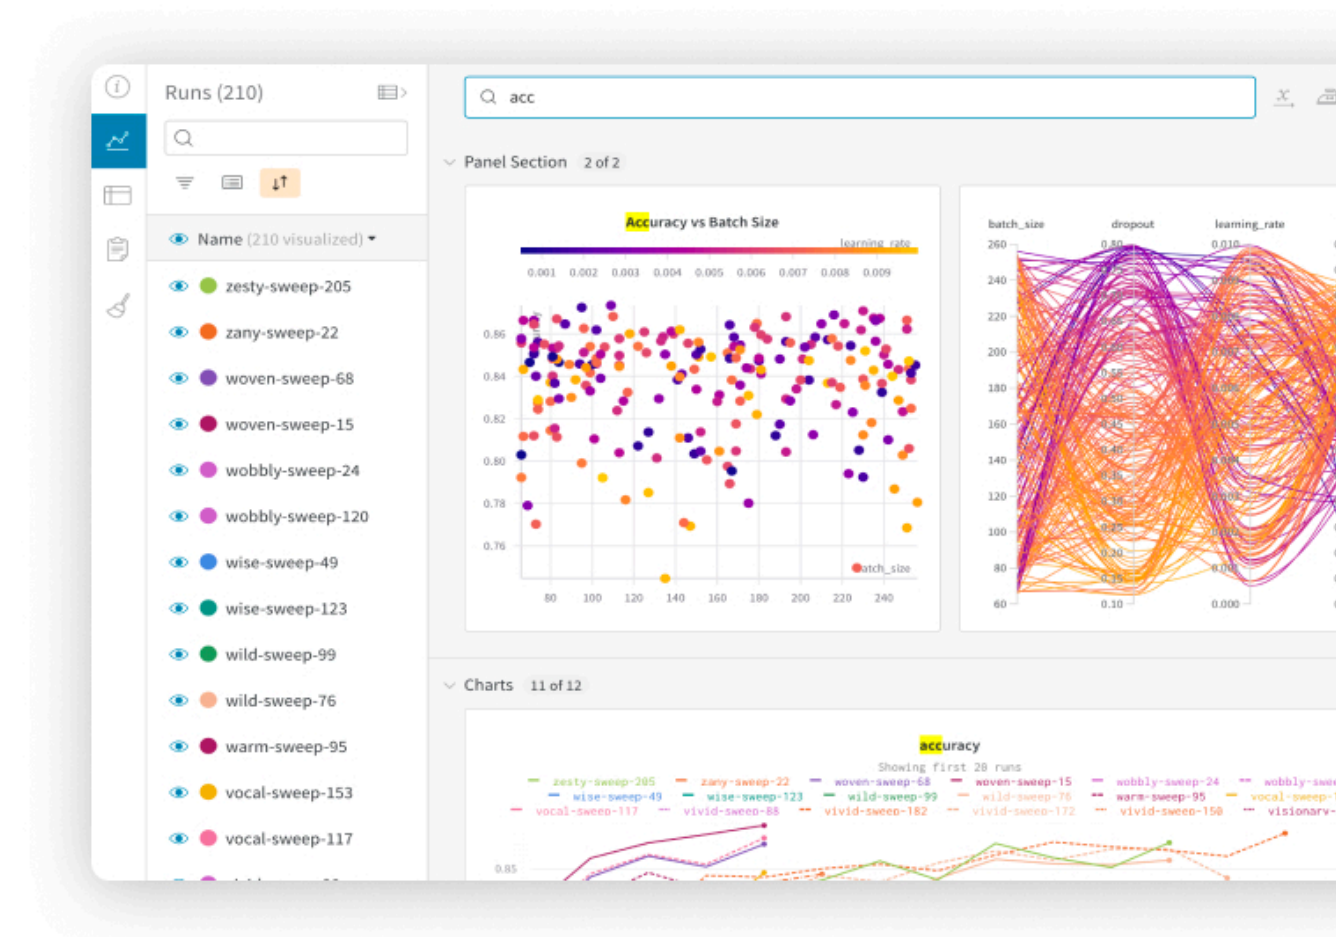
Features of a good experiment tracking tool:

- A central database to store data (configs, metrics, rollouts, etc) from all runs.
- Ability to filter and search by config parameters.
- Ability to tag, modify configs after the run,
- Secure storage, because you do not want to lose your valuable runs :)

**Answer:**



Weights & Biases





# Simple Tracking Demo

00 main\_start.py > ...

```
1  import random
2  import numpy as np
3  import logging
4
5  logger = logging.getLogger("cool_project")
6  logger.setLevel(logging.INFO)
7  logging.basicConfig()
8
9  def single_run(config):
10     # Your crazy RL experiment here
11     for i in range(config['num_steps']):
12         # Craete a dummy learning curve
13         # Dummy exponential decay
14         loss = np.exp(-i / config['decay_rate']) + random.random()
15         # Log the loss
16         logger.info("loss: %f", loss)
17
18
19
20 if __name__ == "__main__":
21     config = {
22         'seed': 42,
23         'num_steps': 100,
24         'decay_rate': 10
25     }
26     single_run(config)
```



# Managing Configs and Log Outputs

**Available tools: Hydra, python-fire**

Features of a good config, log management toolkit:

- Ability to specify complex configs (parent/group configs) through json/yaml.
- Automatically organizes log outputs from runs.
- Ability to specify hyper parameter sweeps without code modifications.
- Ability to override config configurations.

**Answer:**



# Hydra

- A config management tool developed by Facebook.
- Makes it super easy to define hierarchical configs.
- Configs can be overridden using the override keyword, or through CLI
- Automatically manages the run logs into folders.
- Ability to specify hyper-parameter sweeps.

In this example, we configure a server. The server can host multiple websites at the same time.



Output:

```
$ python my_app.py

server:
  site:
    fb:
      domain: facebook.com
    google:
      domain: google.com
  host: localhost
  port: 443
```

# Specifying Sweeps

- Sweeps can be specified using Hydra and run in parallel across:
  - Multiple processes
  - CC array jobs
- All we need to do is write a config file. No code change required!
- The same config can also launch parallel jobs in CC.

```
$ python my_app.py hydra.mode=MULTIRUN db=mysql,postgresql schema=warehouse,support,school
```

```
[2021-01-20 17:25:03,317] [HYDRA] Launching 6 jobs locally
[2021-01-20 17:25:03,318] [HYDRA]      #0 : db=mysql schema=warehouse
[2021-01-20 17:25:03,458] [HYDRA]      #1 : db=mysql schema=support
[2021-01-20 17:25:03,602] [HYDRA]      #2 : db=mysql schema=school
[2021-01-20 17:25:03,755] [HYDRA]      #3 : db=postgresql schema=warehouse
[2021-01-20 17:25:03,895] [HYDRA]      #4 : db=postgresql schema=support
[2021-01-20 17:25:04,040] [HYDRA]      #5 : db=postgresql schema=school
```

# Reproducibility using Singularity Containers

- Replicating exact set of dependencies in CC can be challenging:
  - Packages are often outdated/unavailable.
  - Some dependencies require sudo access.
  - Need to redo installations across all CC clusters

**Answer:**

Singularity Containers



# Singularity

- Singularity is a way to run Docker containers in multi-user clusters such as CC.
- It gives us a way to **package** scientific software and **deploy** them to *different* clusters having the *same* architecture.
- Steps to run Singularity in CC:
  - Start with a Dockerfile (a recipe to build an image)
  - Build a Docker image.
  - Use a different Linux machine to convert Docker image to .sif
  - Copy .sif to CC clusters and issue commands using:

```
$ singularity run -B /home -B /project -B /scratch -B /localscratch:/temp myimage.simg some-program
```



# But, what is a Dockerfile?

Its a recipe to build a Linux image.

```
# System
```

```
FROM nvidia/cuda:11.4.2-cudnn8-devel-ubuntu20.04
```

START WITH BASE IMAGE

```
RUN apt-get update && apt-get install -y \
    ffmpeg git python3-pip vim libglew-dev \
    x11-xserver-utils xvfb \
    && apt-get clean
```

```
RUN pip3 install --upgrade pip
```

```
RUN sh scripts/install-dmlab.sh
```

```
RUN sh scripts/install-atari.sh
```

```
RUN sh scripts/install-minecraft.sh
```

RUN REGULAR LINUX COMMANDS  
TO INSTALL DEPENDENCIES

```
# Agent
```

```
RUN pip3 install jax[cuda11_cudnn82] -f https://storage.googleapis.com/jax-releases/jax\_cuda\_releases.html
```

```
RUN pip3 install jaxlib
```

```
RUN pip3 install tensorflow_probability
```

```
RUN pip3 install optax
```

```
RUN pip3 install tensorflow-cpu
```

```
ENV XLA_PYTHON_CLIENT_MEM_FRACTION 0.8
```

```
(rlt) → Practical Guide Materials docker build -f Dockerfile -t dreamerv3 . --no-cache --platform=linux/arm64
```

BUILD RECIPE INTO A NEW IMAGE



# Useful Resources

- Weights and Biases: <https://docs.wandb.ai/quickstart>
- Hydra: <https://hydra.cc/docs/intro/>
- Submitit: <https://github.com/facebookincubator/submitit>
- Submitit Hydra Launcher Plugin: [https://hydra.cc/docs/plugins/submitit\\_launcher/](https://hydra.cc/docs/plugins/submitit_launcher/)
- Docker getting started: <https://docker-curriculum.com/>
- Singularity installation in Ubuntu 20.04: <https://github.com/apptainer/singularity/issues/5099>
- Singularity in CC: <https://docs.alliancecan.ca/wiki/Singularity>

# Summary

- Having the right set of tools and pipeline can make experimentation easy.
- Various tools could be used to automate various aspects of the experiment pipeline:
  - Experiment Tracking (**Weights and Biases**)
  - Logging and Configs (**Hydra**)
  - Running hyper parameter sweeps (**Submitit plugin**)
- The set of tools often depends on your use-case.